

# CS4532 Concurrent Programming

## Take-Home Lab 1

Fernando T.H.L (210167E)  
Gamage M.S (210176G)

September 4, 2025

## System Information

### CPU

- Model: Intel(R) Xeon(R) Processor @ 2.30GHz
- Vendor/Arch: GenuineIntel / x86-64
- Physical cores: 4
- Threads per core: 1
- Caches:
  - L1i: 128 KiB (4 instances)
  - L1d: 128 KiB (4 instances)
  - L2: 1 MiB (4 instances)
  - L3: 45 MiB (1 instance)

### Memory / NUMA

- Total (kB): 8150140
- NUMA nodes: 1
- THP: madvise
- Swap total (kB): 0

### Operating System

- Distro: Ubuntu 24.04.2 LTS
- Kernel: 6.8.0
- Logical CPUs: 4

### Toolchain

- Compiler: gcc (Ubuntu 13.3.0-6ubuntu2 24.04) 13.3.0
- make: GNU Make 4.3
- glibc: glibc 2.39
- libpthread (NPTL): NPTL 2.39
- Python: 3.12.11
- pandas: 2.3.2
- matplotlib: 3.10.6

## Approach

We implemented a singly linked list supporting:

- **Member**
- **Insert** (unique keys only)
- **Delete**

Three variants were tested:

- Serial (no locks)
- Pthreads + single mutex
- Pthreads + single read-write lock

Initialization:  $n = 1000$  unique keys in  $[0, 2^{16} - 1]$ . Workloads:  $m = 10000$  operations with given fractions, distributed across  $T \in \{1, 2, 4, 8\}$  threads. Timing measures only the  $m$ -operations region, not initialization.

## Experiment Report (Overview Tables)

**Case 1:** n=1000, m=10000, m\_member=0.99, m\_insert=0.005, m\_delete=0.005

Threads	Serial (s)	Mutex (s)	RW-lock (s)
1	0.1815 $\pm$ 0.0000	0.1919 $\pm$ 0.0000	0.2210 $\pm$ 0.0000
2	0.1849 $\pm$ 0.0000	0.2523 $\pm$ 0.0000	0.1579 $\pm$ 0.0000
4	0.1829 $\pm$ 0.0000	0.3796 $\pm$ 0.0000	0.1810 $\pm$ 0.0000
8	0.1780 $\pm$ 0.0000	0.4435 $\pm$ 0.0000	0.2054 $\pm$ 0.0000

Table 1: Summary of results for Case 1.

**Case 2:** n=1000, m=10000, m\_member=0.90, m\_insert=0.05, m\_delete=0.05

Threads	Serial (s)	Mutex (s)	RW-lock (s)
1	1.0639 $\pm$ 0.0000	1.0372 $\pm$ 0.0000	1.0392 $\pm$ 0.0000
2	1.0582 $\pm$ 0.0000	1.2367 $\pm$ 0.0000	1.1793 $\pm$ 0.0000
4	1.0478 $\pm$ 0.0000	1.5626 $\pm$ 0.0000	1.2864 $\pm$ 0.0000
8	1.0130 $\pm$ 0.0000	1.7605 $\pm$ 0.0000	1.4382 $\pm$ 0.0000

Table 2: Summary of results for Case 2.

**Case 3:** n=1000, m=10000, m\_member=0.50, m\_insert=0.25, m\_delete=0.25

Threads	Serial (s)	Mutex (s)	RW-lock (s)
1	4.4753 $\pm$ 0.0000	4.8215 $\pm$ 0.0000	4.5447 $\pm$ 0.0000
2	4.4604 $\pm$ 0.0000	5.8180 $\pm$ 0.0000	7.5133 $\pm$ 0.0000
4	4.4076 $\pm$ 0.0000	6.1148 $\pm$ 0.0000	8.8153 $\pm$ 0.0000
8	4.4076 $\pm$ 0.0000	6.3508 $\pm$ 0.0000	9.0376 $\pm$ 0.0000

Table 3: Summary of results for Case 3.

**Sampling/Confidence** The target of a 5

## Case Analyses with Plots

### Case 1: Read-Heavy Workload

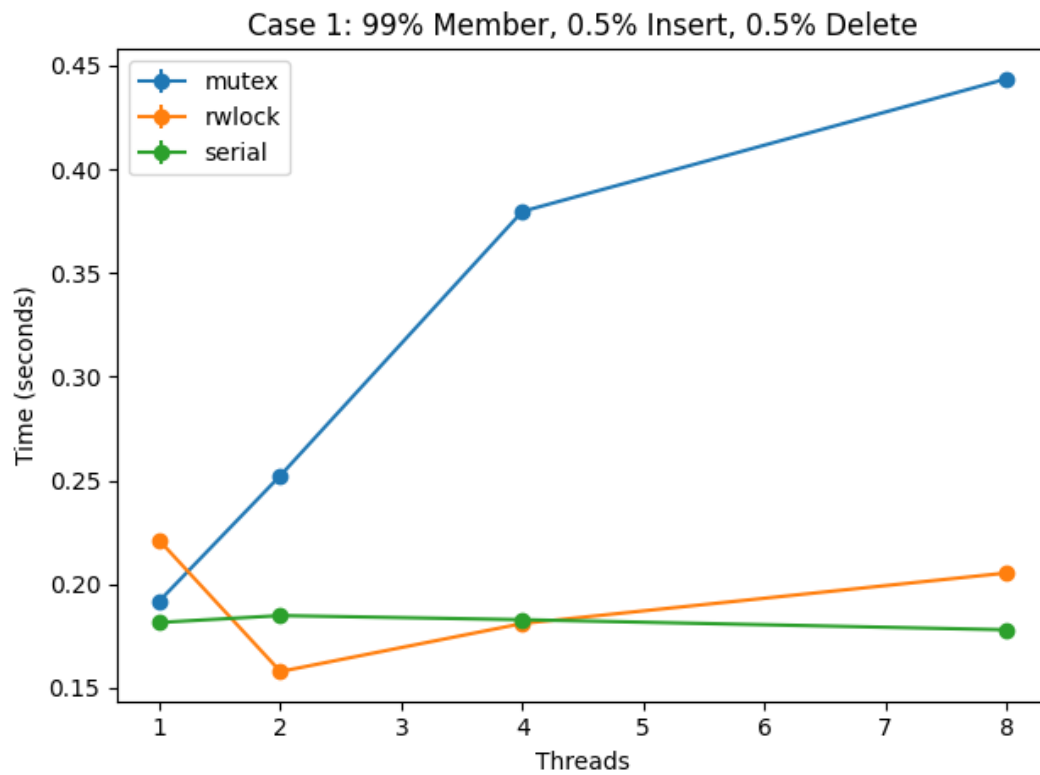


Figure 1: Average time vs. threads for Case 1.

**Analysis** As shown in Table 1 and Figure 1, at 1 thread, serial is fastest (0.1815s) vs mutex (0.1919s) and rw-lock (0.2210s). From 1 to 8 threads, mutex changes by 131.10%. At 8 threads, rw-lock is 2.16x faster than mutex. This workload is read-heavy (99

## Case 2: Balanced Workload

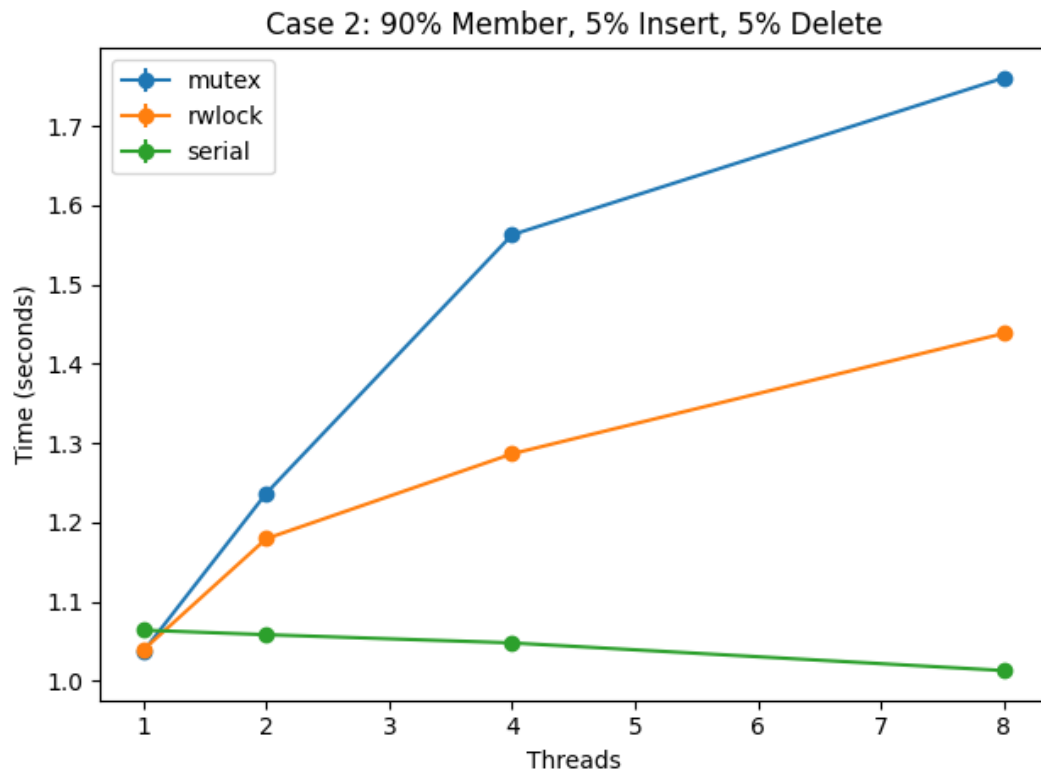


Figure 2: Average time vs. threads for Case 2.

**Analysis** As shown in Table 2 and Figure 2, at 1 thread, serial is fastest (1.0639s) vs mutex (1.0372s) and rw-lock (1.0392s). From 1 to 8 threads, mutex changes by 69.74%. At 8 threads, rw-lock is 1.22x faster than mutex. With a higher write fraction (10

### Case 3: Write-Heavy Workload

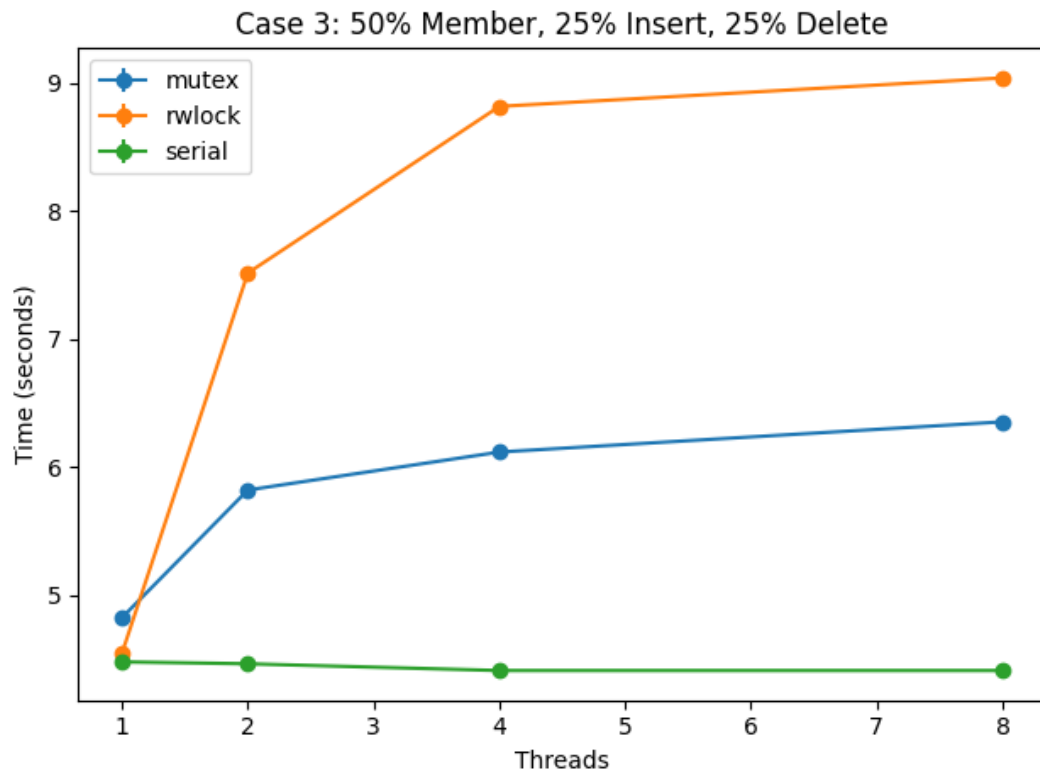


Figure 3: Average time vs. threads for Case 3.

**Analysis** As shown in Table 3 and Figure 3, at 1 thread, serial is fastest (4.4753s) vs mutex (4.8215s) and rw-lock (4.5447s). From 1 to 8 threads, mutex changes by 31.72%. At 8 threads, rw-lock is 0.70x faster than mutex. In this write-heavy scenario (50%

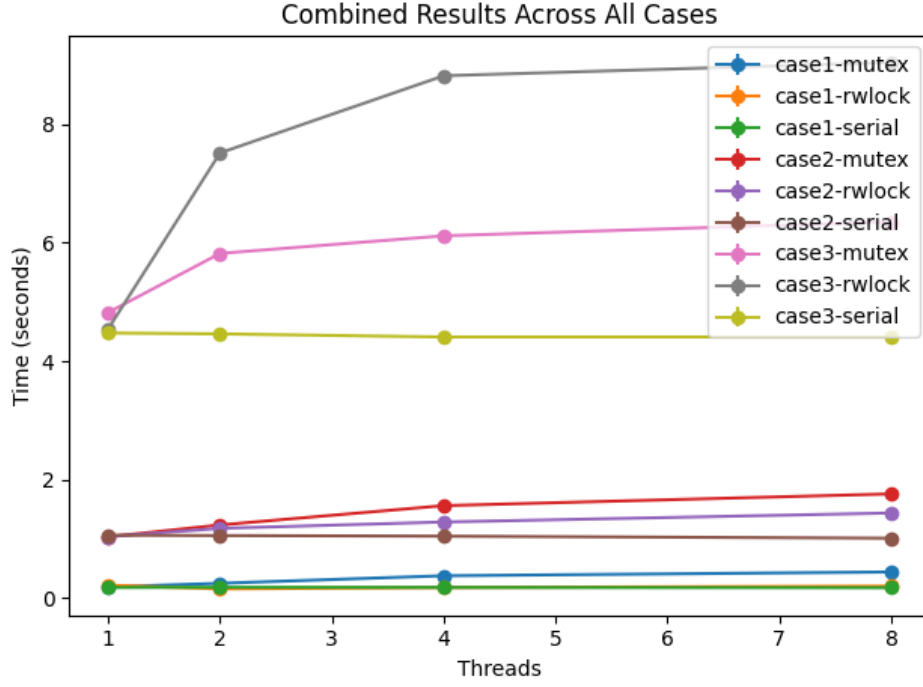


Figure 4: Combined view across all cases and implementations.

## Conclusion

Results align with expectations: the serial baseline dominates at  $T=1$  (no lock overhead). Read-heavy workloads: rwlock outperforms mutex via concurrent readers. Write-heavier workloads: rwlock advantage shrinks; both converge due to writer serialization; parallel versions can underperform serial when contention dominates. Scaling saturates near core count due to contention and scheduling overhead. The  $\pm 5$