

# CS4532 Concurrent Programming

## Take-Home Lab 1

Fernando T.H.L (210167E)  
Gamage M.S (210176G)

September 4, 2025

## System Information

### CPU

- Model: Intel(R) Xeon(R) Processor @ 2.30GHz
- Vendor/Arch: GenuineIntel / x86-64
- Physical cores: 4
- Threads per core: 1
- Caches:
  - L1i: 128 KiB (4 instances)
  - L1d: 128 KiB (4 instances)
  - L2: 1 MiB (4 instances)
  - L3: 45 MiB (1 instance)

### Memory / NUMA

- Total (kB): 8150140
- NUMA nodes: 1
- THP: madvise
- Swap total (kB): 0

### Operating System

- Distro: Ubuntu 24.04.2 LTS
- Kernel: 6.8.0
- Logical CPUs: 4

### Toolchain

- Compiler: gcc (Ubuntu 13.3.0-6ubuntu2 24.04) 13.3.0
- make: GNU Make 4.3
- glibc: glibc 2.39
- libpthread (NPTL): NPTL 2.39
- Python: 3.12.11
- pandas: 2.3.2
- matplotlib: 3.10.6

## Approach

We implemented a singly linked list supporting:

- **Member**
- **Insert** (unique keys only)
- **Delete**

Three variants were tested:

- Serial (no locks)
- Pthreads + single mutex
- Pthreads + single read-write lock

Initialization:  $n = 1000$  unique keys in  $[0, 2^{16} - 1]$ . Workloads:  $m = 10000$  operations with given fractions, distributed across  $T \in \{1, 2, 4, 8\}$  threads. Timing measures only the  $m$ -operations region, not initialization.

## Experiment Report (Overview Tables)

**Case 1:** n=1000, m=10000, m\_member=0.99, m\_insert=0.005, m\_delete=0.005

Threads	Serial ( $\mu$ s)	Mutex ( $\mu$ s)	RW-lock ( $\mu$ s)
1	9252.00 $\pm$ 127.00	10112.00 $\pm$ 459.00	10930.00 $\pm$ 243.00
2	9242.00 $\pm$ 75.00	32093.00 $\pm$ 5094.00	7674.00 $\pm$ 250.00
4	8955.00 $\pm$ 288.00	30534.00 $\pm$ 2223.00	5229.00 $\pm$ 446.00
8	8832.00 $\pm$ 212.00	29664.00 $\pm$ 1443.00	6179.00 $\pm$ 1148.00

Table 1: Summary of results for Case 1.

**Case 2:** n=1000, m=10000, m\_member=0.90, m\_insert=0.05, m\_delete=0.05

Threads	Serial ( $\mu$ s)	Mutex ( $\mu$ s)	RW-lock ( $\mu$ s)
1	17615.00 $\pm$ 825.00	18042.00 $\pm$ 462.00	29770.00 $\pm$ 533.00
2	17107.00 $\pm$ 182.00	39649.00 $\pm$ 5609.00	30232.00 $\pm$ 908.00
4	17372.00 $\pm$ 94.00	42669.00 $\pm$ 2007.00	22977.00 $\pm$ 823.00
8	17664.00 $\pm$ 173.00	48538.00 $\pm$ 469.00	22568.00 $\pm$ 5037.00

Table 2: Summary of results for Case 2.

**Case 3:** n=1000, m=10000, m\_member=0.50, m\_insert=0.25, m\_delete=0.25

Threads	Serial ( $\mu$ s)	Mutex ( $\mu$ s)	RW-lock ( $\mu$ s)
1	59140.00 $\pm$ 540.00	61343.00 $\pm$ 512.00	68963.00 $\pm$ 2772.00
2	58727.00 $\pm$ 661.00	85704.00 $\pm$ 6087.00	107904.00 $\pm$ 4267.00
4	58024.00 $\pm$ 773.00	103935.00 $\pm$ 6010.00	120342.00 $\pm$ 4358.00
8	57798.00 $\pm$ 1239.00	115380.00 $\pm$ 2375.00	121688.00 $\pm$ 5765.00

Table 3: Summary of results for Case 3.

**Sampling/Confidence** For Case 1, the worst relative CI was 16.29For Case 2, the worst relative CI was 19.56For Case 3, the worst relative CI was 6.23The target of a 5

## Case Analyses with Plots

### Case 1: Read-Heavy Workload




Figure 1: Average time vs. threads for Case 1.

**Analysis** As shown in Table 1 and Figure 1, at 1 thread, serial is fastest (9252.00μs) vs mutex (10112.00μs) and rw-lock (10930.00μs). From 1 to 8 threads, mutex changes by 193.35At 8 threads, rw-lock is 4.80x faster than mutex. This workload is read-heavy (99

## Case 2: Balanced Workload




Figure 2: Average time vs. threads for Case 2.

**Analysis** As shown in Table 2 and Figure 2, at 1 thread, serial is fastest (17615.00μs) vs mutex (18042.00μs) and rw-lock (29770.00μs). From 1 to 8 threads, mutex changes by 169.03%. At 8 threads, rw-lock is 2.15x faster than mutex. With a higher write fraction (10

### Case 3: Write-Heavy Workload




Figure 3: Average time vs. threads for Case 3.

**Analysis** As shown in Table 3 and Figure 3, at 1 thread, serial is fastest (59140.00μs) vs mutex (61343.00μs) and rw-lock (68963.00μs). From 1 to 8 threads, mutex changes by 88.09%. At 8 threads, rw-lock is 0.95x faster than mutex. In this write-heavy scenario (50




Figure 4: Combined view across all cases and implementations.

## Conclusion

Results align with expectations: the serial baseline dominates at  $T=1$  (no lock overhead). Read-heavy workloads: rwlock outperforms mutex via concurrent readers. Write-heavier workloads: rwlock advantage shrinks; both converge due to writer serialization; parallel versions can underperform serial when contention dominates. Scaling saturates near core count due to contention and scheduling overhead. The  $\pm 5$